




AI X MARKETING · EU

ChatGPT text in the EU is getting an invisible watermark.

What it can show, what it can't, and what I'd
change in your workflow.




swipe →




01 · The news

OpenAI announced it on 5 October.

Over the coming weeks, eligible ChatGPT and Codex text made in the EU gets an invisible watermark. It's OpenAI's answer to the [EU AI Act](#), which wants AI text to be machine-detectable.




swipe →



02 · How it works

It hides in word choice, not on screen.

The method, called textGrain, nudges which words the model picks. A detector looks for that statistical pattern. Your text looks exactly the same.





03 · How reliable

~80%

of 200-token passages detected.

About 95% at 400 tokens, at a 1% false-positive target. Shorter text and heavily edited text are much harder to detect. OpenAI says so itself.

Source: OpenAI, 5 Oct 2026



What a watermark can't tell anyone

CAN'T TELL

- ◆ Who wrote or prompted it
- ◆ How much a human changed
- ◆ Whether it's true
- ◆ Who owns it

CAN SIGNAL


- ◆ That an OpenAI model made or processed part of it
- ◆ Only if the passage is long enough



05 · Who can check

Not your client. Not yet.

The detector isn't public. At launch, only approved researchers and expert organisations can apply. In the API, watermarking is opt-in and **off by default**.



swipe →



What I'd do this week

- 01 Treat AI output as a draft. Editing for voice and facts was always the job.
 - 02 Check client and creator contracts for AI-disclosure clauses.
 - 03 Decide your own disclosure rule before a client asks.
 - 04 Ignore anyone selling 'AI detector' screenshots as proof.
-



SAVE FOR LATER

Save this before your next **client call.**

Full breakdown, sources and a disclosure
template on Build with Vish.

YOUR TURN

Would you tell a client a draft was AI-assisted? Yes or
no?

Save · Follow for AI x marketing

